

The 7-Point SOP Validation Framework

A named methodology for writing SOPs someone can actually execute under pressure, and confirming they work before you rely on them. Prepared 2026-08-13. Updated 2026-08-13 with patterns confirmed across a 13-SOP Amazon marketplace operations validation series — see the companion Validation Log for the SOP-by-SOP findings.

Writing Best Practices

- 1. Scope it to one trigger, one owner, one outcome.** The moment an SOP tries to cover two different situations (e.g., "customer complaints and returns"), it becomes two SOPs wearing one name. Split them, even if they share a department.
- 2. Write for the person with the least context, not the most.** Assume the reader has never seen this exact situation before. "Check the account health page" is a description; "log into Seller Central → Account Health → note the exact policy cited in the case ID" is an instruction. If a step requires unstated tribal knowledge to execute, it isn't finished.
- 3. Every decision point must be a testable condition, not a judgment call.** "If the issue seems serious" is not usable. "If claim value >\$500" is. If you can't write the IF as a number, a named category, or a yes/no fact, the SOP is hiding a decision that needs its own escalation path instead.
- 4. One imperative action per numbered step.** Not "review the report and update the tracker if needed" — that's two steps disguised as one, and "if needed" is an unowned decision point again.
- 5. Name the role and the tool at every handoff, not just the action.** RACI tables catch this — if a step has no clear "R," it won't happen reliably during an actual incident, only during the tabletop read-through.
- 6. Numbers over adverbs.** "Promptly," "as soon as possible," and "regularly" are not SLAs. Tie a time bound to a severity tier for every commitment — it's the difference between an aspiration and something you can audit against later.
- 7. Separate reactive from scheduled structurally, not just in content.** A Response SOP (triggered by an unpredictable event — a suspension, a counterfeit listing) needs a decision tree and severity tiers. A Process SOP (triggered by the calendar — a weekly inventory review, a monthly ad spend review) needs a cadence and a linear procedure. Forcing one format onto both makes the reactive ones too rigid and the scheduled ones full of unnecessary branching.

Validation Methodology

- 1. Blind execution test.** Hand the SOP to someone who didn't write it and wasn't involved in the underlying process, and have them execute it literally, step by step, without asking you anything. Count how many times they have to guess or ask for clarification — the target is zero. Every clarification they need is a gap in the document, not a gap in their competence.

- 2. Audit against real past incidents.** Pull 3–5 actual cases that already happened and manually trace them through the SOP's decision tree. If the tree doesn't land on the outcome that actually happened, the branching logic is wrong or incomplete.
- 3. Check for orphaned branches.** Every “IF” in the decision tree needs a resolution path that's actually written out below it. This is the single most common SOP defect — a trigger condition gets listed, but no one ever wrote what to do about it.
- 4. Independent-execution consistency check.** Have two different people execute the same trigger scenario independently, without conferring. If they produce different actions or different outcomes, the SOP has ambiguity somewhere, even if both people “followed the steps.”
- 5. Time it against the stated SLA.** If the document says “2-hour first response” but a live run takes five, either the SLA is unrealistic or a step is missing — for example, no one told the executor where to find the case ID quickly.
- 6. Escalation path stress test.** Actually contact the named escalation role during a drill and confirm they know they're the point person for that condition. A correct org chart on paper and a functioning escalation path in practice are not the same thing.
- 7. Track deviation rate, not compliance rate.** If people are consistently doing something different from what's written, that's not a training problem to fix by enforcement — it's a signal the SOP doesn't match operational reality and needs revision. This is what a Revision History and a fixed review cadence are actually for, not just version-stamping.

Confirmed Patterns — From the 13-SOP Validation Series

Running this checklist against all 13 SOPs in the framework (BP-001 through CM-001) surfaced the same handful of defects independently, across unrelated authors and functions. These four are no longer theoretical — treat them as standing rules for every future SOP and every future revision.

- 1. Every RACI row needs exactly one Accountable role — no more, no fewer.** Zero Accountables (3 confirmed instances, across AH-001 and FR-001) means nobody owns the outcome. Two Accountables (5 confirmed instances, across PA-001, PD-001, AD-001, RR-001, and CM-001) means accountability is diffused — each party can assume the other is driving. The dual-Accountable failure clustered specifically on steps where a client and an internal lead jointly participate: a client review call, calendar planning, budget approval, product intake. Treat any RACI row with that shape as a place to double-check before it ships.
- 2. Every role named in a procedure, trigger, or escalation clause must appear in that SOP's own RACI table.** Confirmed 4 times (IM-001, LF-001, AH-001, PC-001), always the same shape: an escalation or notification step names a role — most often Account Manager — that was never actually defined as a stakeholder in Section 3. If a role can receive an escalation or a hand-off, it needs a row, even if that row is mostly “I” (Informed).
- 3. Every escalation needs a resolution bound stated separately from its response bound.** A trigger (“escalates within 24 hours”) tells you when to hand something off; it says nothing about how fast the receiving party must act once they have it. This gap appeared in nearly every SOP in the series until fixed. Treat “escalates to X” as an incomplete sentence until it also states what X does and by when.

4. Cross-SOP dependencies must be re-verified after either side changes, not just checked once. LF-001 and FR-001 shared a trigger condition with no cross-reference at all — a real dollar-risk gap, since a receiving shortage could be "resolved" on the logistics side while its 60-day Amazon reimbursement window quietly expired. AD-001 and CM-001 originally had no relationship; fixing AD-001 to reference CM-001 introduced a new contradiction with CM-001's own wording, which only surfaced when CM-001 was independently validated afterward. Whenever a fix to one SOP references another, check the referenced SOP's actual current wording, not just that the reference exists — a fix made in isolation can silently break consistency with a document it now depends on.

Benchmark Reference

Waymo's Emergency Response Guide and Law Enforcement Interaction Protocol (see Waymo_SOP_Reference folder) is a useful external benchmark for items 1 and 3 above — it's written for first responders who have never touched a Waymo vehicle before and have seconds to act, so every branch is exhaustive and nothing assumes prior context. That is the bar a Response-variant SOP should be held to.